

2.2 Bottleneck Identification

Answer:

The system faces five potential bottlenecks:

1. Database queries

- Redis may become overloaded by frequent feature lookups.
- TimescaleDB can slow down due to large scans of time-series data.
- ClickHouse may struggle with expensive ad-hoc analytics queries.

Optimization: Use pipelined queries and sharding in Redis, partition and rollups in TimescaleDB, and materialized views in ClickHouse.

2. API rate limits

- Hyperliquid APIs may reject orders if too many requests are sent at once.
- Third-party APIs for sentiment or economic data can also throttle requests.

Optimization: Apply token-bucket rate limiting, use idempotency keys, batch operations, and introduce caching with short TTLs to reduce redundant requests.

3. Network latency

- Cross-region traffic and repeated serialization can slow down data delivery.

Optimization: Co-locate latency-sensitive components, use long-lived WebSocket connections, and compress messages using binary formats.

4. Computational overhead

- Stream processors may waste time recomputing windows or parsing JSON.
- Agents can slow down due to heavy inference or noisy signals.

Optimization: Use incremental window updates, vectorized math (NumPy/Numba), batch inference, and debounce agent decisions.

5. Concurrency issues

- Multiple agents could race to trade the same symbol, or the router could become a bottleneck.

Optimization: Use a Redis “blackboard” for intents, enforce a central risk gate for final approval, and design the order router as stateless replicas with shared idempotency storage.

2.3 Data Pipeline Design

Answer:

1. Real-time price feeds

- Hyperliquid WebSocket → normalize messages → event bus → stream processors → Redis feature store → agents.
- This design ensures sub-500 ms latency because all steps are in-memory and event-driven.

2. Historical data

- Subgraph/indexing providers (e.g., Envio/Goldsky) are used to fetch on-chain history.
- Results are stored in TimescaleDB (time-series queries) and ClickHouse (analytics).
- Cached in Redis for frequently requested queries.

3. Extensibility

- New data sources (e.g., Twitter sentiment) plug in through a connector interface (`connect()` , `normalize()` , `emit()`).
- All connectors output standardized messages into the event bus, so the core system does not need changes.

4. Data freshness vs. consistency

- For trading, prioritize freshness: agents act on the latest available data, even if slightly inconsistent.
- For reporting, prioritize consistency: use reconciled subgraph data for audits.

- Features include timestamps; agents enforce staleness rules (e.g., price <500 ms old, sentiment <5 minutes).
-

2.4 Multi-Agent State Management

Answer:

- **Sharing information without races:** Agents publish their latest trading intent to a shared Redis "blackboard" (`intent:BTC = {agent, side, size, ts}`), which others can read. This avoids direct overwrites.
- **Preventing conflicting orders:** All intents pass through a central risk engine. It enforces exposure, leverage, and rate limits before forwarding approved intents to the order router. The router uses idempotency keys to avoid duplicate orders.
- **Maintaining audit trails:** Every step (intent creation, approval, submission, fills) is logged as an append-only event in TimescaleDB for operations and ClickHouse for analytics. Each entry includes IDs and timestamps to enable full replay.
- **Handling crashes:**
 - Agents send heartbeats; if one fails, the risk engine cancels its pending intents and the router cleans up orders.
 - Routers are stateless replicas, so they can restart and replay safely without duplicates.
 - Ingestion services can replay missed events from the event bus.

Summary: Redis provides fast context sharing, the risk engine ensures safe execution, audit trails guarantee traceability, and heartbeats with idempotency provide resilience.
